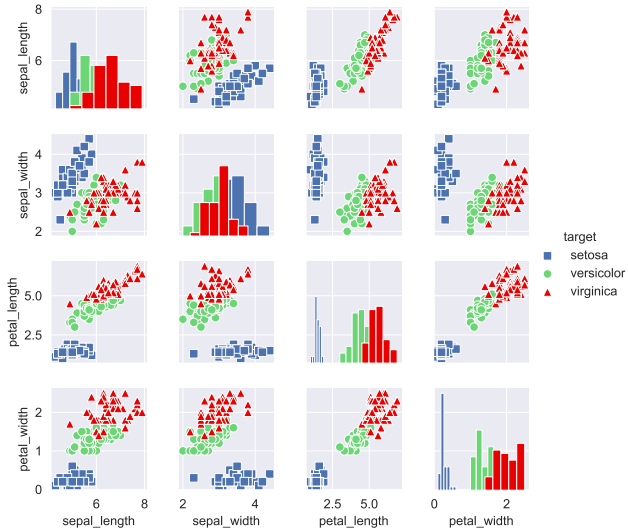# Machine Learning

Lecture 2: k-Nearest Neighbors

Prof. Dr. Stephan Günnemann
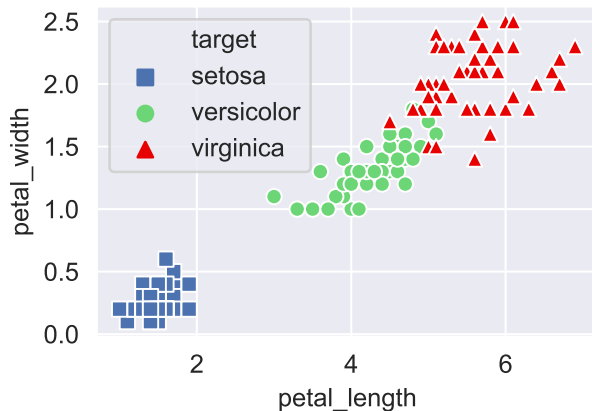
Data Analytics and Machine Learning
Technical University of Munich

Winter term 2021/2022

Iris dataset

Data Analytics and
Machine Learning

Iris dataset: 2 features

How do we intuitively label new samples by hand?
Look at the *surrounding* points. Do as your neighbor does.

1-NN algorithm

Given a training dataset $\mathcal{D} = \{(\boldsymbol{x}_i, y_i)\}_{i=1}^N$
where $\boldsymbol{x}_i \in \mathbb{R}^D$ are features and $y_i \in \{1, \ldots, C\}$ are class labels

To classify new observations:

- define a distance measure (e.g. Euclidean distance)

- compute the nearest neighbor for all new data points

- and label them with the label of their nearest neighbor

This works for both *classification* and *regression*.

1-NN

This corresponds to a Voronoi tesselation.
And results in poor generalization...

k-Nearest Neighbor classification

More *robust* against errors in the training set:

Look at multiple nearest neighbors and pick the majority label.

k-Nearest Neighbor classification

More *robust* against errors in the training set:

Look at multiple nearest neighbors and pick the majority label.

Let $\mathcal{N}_k(\boldsymbol{x})$ be the k nearest neighbors of a vector $\boldsymbol{x}$, then in classification tasks:

$$p(y = c \mid \boldsymbol{x}, k) = \frac{1}{k} \sum_{i \in \mathcal{N}_k(\boldsymbol{x})} \mathbb{I}(y_i = c),$$

$$\hat{y} = \arg\max_c p(y = c \mid \boldsymbol{x}, k)$$

with the *indicator variable* $\mathbb{I}(e)$ is defined as:

$$\mathbb{I}(e) = \begin{cases} 1 \text{ if } e \text{ is true} \\ 0 \text{ if } e \text{ is false.} \end{cases}$$

k-Nearest Neighbor classification

More *robust* against errors in the training set:

Look at multiple nearest neighbors and pick the majority label.

Let $\mathcal{N}_k(\boldsymbol{x})$ be the k nearest neighbors of a vector $\boldsymbol{x}$, then in classification tasks:

$$p(y = c \mid \boldsymbol{x}, k) = \frac{1}{k} \sum_{i \in \mathcal{N}_k(\boldsymbol{x})} \mathbb{I}(y_i = c),$$

$$\hat{y} = \arg \max_c p(y = c \mid \boldsymbol{x}, k)$$

with the *indicator variable* $\mathbb{I}(e)$ is defined as:

$$\mathbb{I}(e) = \begin{cases} 1 \text{ if } e \text{ is true} \\ 0 \text{ if } e \text{ is false.} \end{cases}$$

i.e., the vector will be labeled by the mode of its neighbors' labels.

k-Nearest Neighbor classification: weighted

Look at multiple nearest neighbors and pick the weighted majority label.

k-Nearest Neighbor classification: weighted

Look at multiple nearest neighbors and pick the weighted majority label. The weight is inversely proportional to the distance.

Let $\mathcal{N}_k(\boldsymbol{x})$ be the k nearest neighbors of a vector $\boldsymbol{x}$, then in classification tasks:

$$p(y = c \mid \boldsymbol{x}, k) = \frac{1}{Z} \sum_{i \in \mathcal{N}_k(\boldsymbol{x})} \frac{1}{\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)} \mathbb{I}(y_i = c),$$

$$\hat{y} = \arg \max_c p(y = c \mid \boldsymbol{x}, k)$$

with $Z = \sum_{i \in \mathcal{N}_k(\boldsymbol{x})} \frac{1}{\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)}$ the normalization constant and $\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)$ being a distance measure between $\boldsymbol{x}$ and $\boldsymbol{x}_i$.

k-Nearest-Neighbor regression

Regression is similar:

Let $\mathcal{N}_k(\boldsymbol{x})$ be the k nearest neighbors of a vector $\boldsymbol{x}$, then for regression:

$$\hat{y} = \frac{1}{Z} \sum_{i \in \mathcal{N}_k(\boldsymbol{x})} \frac{1}{\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)} \, y_i,$$

with $Z = \sum_{i \in \mathcal{N}_k(x)} \frac{1}{\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)}$ the normalization constant and $\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)$ being a distance measure between $\boldsymbol{x}$ and $\boldsymbol{x}_i$,

k-Nearest-Neighbor regression

Regression is similar:

Let $\mathcal{N}_k(\boldsymbol{x})$ be the k nearest neighbors of a vector $\boldsymbol{x}$, then for regression:

$$\hat{y} = \frac{1}{Z} \sum_{i \in \mathcal{N}_k(\boldsymbol{x})} \frac{1}{\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)} \, y_i,$$

with $Z = \sum_{i \in \mathcal{N}_k(x)} \frac{1}{\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)}$ the normalization constant and $\mathrm{d}(\boldsymbol{x}, \boldsymbol{x}_i)$ being a distance measure between $\boldsymbol{x}$ and $\boldsymbol{x}_i$,

i.e., the vector will be labeled by a weighted mean of its neighbors' values.

Note: y_i is a real number here (rather then categorical label).

Data Analytics and
Machine Learning

3-NN

So, how many neighbors are best?

Data Analytics and
Machine Learning

1-NN

Compare the decision boundaries of 1-NN and 3-NN

Data Analytics and
Machine Learning

Choosing k

Goal is generalization: pick k (called a *hyper-parameter*) that performs best[1] on unseen (future) data.

Unfortunately, no access to future data, so split the dataset $\mathcal{D}$:

Training set $\mathcal{D}_T$	Validation set $\mathcal{D}_V$

| Learning set $\mathcal{D}_L$	Test set $\mathcal{D}_t$

$$\mathcal{D}$$

Hyper-parameter tuning procedure

- Learn the model using the training set
- Evaluate performance with different k on the *validation set* picking the best k
- Report final performance on the test set.[2]

---
[1]In terms of some predefined metric, e.g., accuracy
[2]Good data science practices: See slides on Decision Trees

Using validation set to choose k

We choose $k = 7$.

Measuring classification performance

How can we assess the performance of a (binary) classification algorithm?

$\Rightarrow$ Confusion table

Predicted	True condition $y = 1$	$y = 0$
$y = 1$	TP	FP
$y = 0$	FN	TN

$\left.\begin{array}{l} TP = \text{true positive} \\ TN = \text{true negative} \end{array}\right\}$ correct predictions

$\left.\begin{array}{l} FP = \text{false positive} \\ FN = \text{false negative} \end{array}\right\}$ wrong predictions

false negatives   true negatives

true positives   false positives

Image source: `https://en.wikipedia.org/wiki/Precision_and_recall`

Measuring classification performance

Accuracy:
$$\text{acc} = \frac{TP + TN}{TP + TN + FP + FN}$$

Precision:
$$\text{prec} = \frac{TP}{TP + FP}$$

Sensitivty/Recall:
$$\text{rec} = \frac{TP}{TP + FN}$$

Specificity:
$$\text{tnr} = \frac{TN}{FP + TN}$$

False Negative Rate:
$$\text{fnr} = \frac{FN}{TP + FN}$$

False Positive Rate:
$$\text{fpr} = \frac{FP}{FP + TN}$$

F1 Score:
$$f1 = \frac{2 \cdot \text{prec} \cdot \text{rec}}{\text{prec} + \text{rec}}$$

$\Rightarrow$ Trade-off between precision and recall: increasing one (most often) leads to decreasing the other

General note: Be careful when you have imbalanced classes!

Distance measures

- K-NN can be used with various distance measures $\rightarrow$ highly flexible
- Euclidean distance (L_2 norm): $\sqrt{\sum_i (u_i - v_i)^2}$

Distance measures

- K-NN can be used with various distance measures $\rightarrow$ highly flexible
- Euclidean distance (L_2 norm): $\sqrt{\sum_i (u_i - v_i)^2}$

- L_1 norm: $\sum_i |u_i - v_i|$

Distance measures

- K-NN can be used with various distance measures $\rightarrow$ highly flexible
- Euclidean distance (L_2 norm): $\sqrt{\sum_i (u_i - v_i)^2}$

- L_1 norm: $\sum_i |u_i - v_i|$
- L_∞ norm: $\max_i |u_i - v_i|$

Distance measures

- K-NN can be used with various distance measures $\rightarrow$ highly flexible
- Euclidean distance (L_2 norm): $\sqrt{\sum_i (u_i - v_i)^2}$

- L_1 norm: $\sum_i |u_i - v_i|$
- L_∞ norm: $\max_i |u_i - v_i|$
- Angle:

$$\cos \alpha = \frac{\boldsymbol{u}^T \boldsymbol{v}}{\|\boldsymbol{u}\|\|\boldsymbol{v}\|}$$

Distance measures

- K-NN can be used with various distance measures $\rightarrow$ highly flexible
- Euclidean distance (L_2 norm): $\sqrt{\sum_i (u_i - v_i)^2}$

- L_1 norm: $\sum_i |u_i - v_i|$
- L_∞ norm: $\max_i |u_i - v_i|$
- Angle:

$$\cos \alpha = \frac{\boldsymbol{u}^T \boldsymbol{v}}{\|\boldsymbol{u}\|\|\boldsymbol{v}\|}$$

- *Mahalanobis distance* ($\boldsymbol{\Sigma}$ is *positive (semi) definite* and *symmetric*):

$$\sqrt{(\boldsymbol{u} - \boldsymbol{v})^T \boldsymbol{\Sigma}^{-1} (\boldsymbol{u} - \boldsymbol{v})}$$

- Hamming distance, Edit distance, . . .

Data Analytics and
Machine Learning

Scaling issues

The same old example but one of our features is in the order of meters, the other in the order of centimeters. ($k = 1$)

Data Analytics and
Machine Learning

Circumventing scaling issues

- Data *standardization*
  Scale each feature to zero mean and unit variance.

  $$x_{i,\text{std}} = \frac{x_i - \mu_i}{\sigma_i}$$

  (This is a standard procedure in machine learning. Many models are
  sensitive to differences in scale.)

Circumventing scaling issues

- Data *standardization*
  Scale each feature to zero mean and unit variance.

  $$x_{i,\text{std}} = \frac{x_i - \mu_i}{\sigma_i}$$

  (This is a standard procedure in machine learning. Many models are sensitive to differences in scale.)

- Use the Mahalanobis distance.

  $$\text{mahalanobis}(\boldsymbol{x}_1, \boldsymbol{x}_2) = \sqrt{(\boldsymbol{x}_1 - \boldsymbol{x}_2)^T \boldsymbol{\Sigma}^{-1} (\boldsymbol{x}_1 - \boldsymbol{x}_2)}$$

  $\boldsymbol{\Sigma} = \begin{bmatrix} \sigma_1^2 & 0 & 0 \\ 0 & \cdots & 0 \\ 0 & 0 & \sigma_n^2 \end{bmatrix}$ is equal to Euclidean distance on normalized data

Data Analytics and
Machine Learning

The curse of dimensionality

Given a discrete one-dimensional input space
$x \in \{1, 2, \ldots, 10\}$

For $N = 20$ uniformly distributed samples the
data covers 100% of the input space.

The curse of dimensionality

Given a discrete one-dimensional input space $x \in \{1, 2, \ldots, 10\}$

For $N = 20$ uniformly distributed samples the data covers 100% of the input space.

Add a second dimension (now $\boldsymbol{x} \in \{1, \ldots, 10\}^2$) and your data only covers 18% of the input space.

Data Analytics and
Machine Learning

The curse of dimensionality

Given a discrete one-dimensional input space
$x \in \{1, 2, \ldots, 10\}$

For $N = 20$ uniformly distributed samples the
data covers 100% of the input space.

Add a second dimension (now
$\boldsymbol{x} \in \{1, \ldots, 10\}^2$) and your data only covers
18% of the input space.

Once you add a third dimension you only cover
2%.

Data Analytics and
Machine Learning

The curse of dimensionality

- The nearest neighbor will now be pretty far away..
- N has to grow exponentially with the number of features. Consider this when using k-NN on high-dimensional data.

Practical considerations

Expensive: memory and naive inference are both $O(N)$:

we need to store the entire training data and compare with all training instances to find the nearest neighbor

Solution: use tree-based search structures (e.g. k-d tree) for efficient (approximate) NN [3]

---

[3]At the expense of an additional computation performed only once

What we learned

- k-NN Algorithm
- Train-validation-test split
- Measuring classification performance
- Distance metrics
- Curse of dimensionality

Reading material

Main reading

- "Machine Learning: A Probabilistic Perspective" by Murphy [ch. 1.4.1 - 1.4.3]

Extra reading

- "Bayesian Reasoning and Machine Learning" by Barber [ch. 14]

Slides adapted from previous versions by W. Koepp & D. Korhammer

Data Analytics and
Machine Learning